

Multiagent Learning in an Real Time Strategy Environment

Bachelor's Thesis of

Christian Burmeister

at the Department of Informatics
Institute for Anthropomatics and Robotics (IAR)
Autonomous Learning Robots (ALR)

Reviewer:	Prof. A
Second reviewer:	Prof. B
Advisor:	M.Sc. C
Second advisor:	M.Sc. D

xx. Month 2021 – xx. Month 2021

I declare that I have developed and written the enclosed thesis completely by myself, and have not used sources or means without declaration in the text.

PLACE, DATE

.....
(Christian Burmeister)

Abstract

This thesis explores a multitude of approaches for crafting an AI for a real-time strategy video game that uses learning systems to provide an adaptive system. This would allow for a sophisticated player experience with little external input during production and after shipment. Given that RTS games construct a complex and rich simulation environment for single-agent and multi-agent systems they allow advancements in AI research to be benchmarked virtually and flexibly in an efficient manner. It will become apparent that a lot of the challenges can be dissected into loosely coupled systems which can be solved individually by different AI and learning methods. This paper focuses on the usage of learning hierarchical multi-agent systems for unit control and tactical maneuvering.

First of all, it is going into more depth about why this research is interesting for the perspective of AI-Research and video game production and especially in comparison with classical video game AI approaches. Then it will go through theoretical fundamentals which are needed to be understood for this thesis and talks about how you would dissect the AI into different parts. Afterwards it will go over what is needed to solve the specific problem for the included project with different solutions and approaches, while focusing on multi-agent systems. The project will entail the implementation of a part of the aforementioned dissected construct with the Godot Engine and a simplified prototype of a RTS game. In closing, the thesis will go over what still needs to be done for a complete AI suite and what aspects can still be improved upon.

(related works chapter???)

Zusammenfassung

Contents

Abstract	i
Zusammenfassung	ii
1. Motivation	2
1.1. Videogames and AI Research	2
1.1.1. Real-Time Strategy Games for AI	2
1.2. Multiagent Systems	2
1.3. Spacing and indentation	3
1.4. Example: Citation	3
1.5. Example: Figures	3
1.6. Example: Tables	3
1.7. Example: Formula	3
2. Information to sort	5
2.1. Books	5
2.2. MARL - A Comprehensive Survey of Multiagent Reinforcement Learning - 2008	5
2.3. Artificial Intelligence - A modern Approach	6
2.3.1. Agents and Environments	6
2.3.2. Rational Agent	6
2.3.3. Nature of Environments	7
2.3.4. Structure of Agents	8
2.3.5. Multiagent Planning	9
2.3.6. Game Theory	11
2.3.7. Mechanism Design for Multiple Agents	11
2.3.8. Adversarial Search	11
2.3.9. Probabilistic Reasoning over Time	11
2.3.10. Reinforcement Learning	11
2.3.11. Planning Uncertain Movements (Potential Fields)	11
2.4. RTSAI - A Review of Real-Time Strategy Game AI - 2014	11
2.4.1. Tactical Decision making	11
2.4.2. Strategic Decision making	12
2.4.3. Open Research Areas	13
2.4.4. AI Techniques that can be used	14
2.5. Ant Colony Optimization	14
2.5.1. Wikipedia Article	14
2.6. UNSORTED	15

2.7.	References & Papers	15
2.7.1.	Ant Colony Optimization (ACO)	15
2.7.2.	Multi Agent Reinforcement Learning (MARL)	15
2.7.3.	Multiagent Systems (MAS)	16
2.7.4.	RTS AI	16
2.7.5.	RTS Learning	17
3.	Fundamentals	18
4.	Problem and Approaches	20
4.1.	Definition of the Problem domain	20
4.2.	Problems and Aspects for an RTS AI	20
4.2.1.	Tactical Decision Making	20
4.2.2.	Strategic Decision Making	20
5.	Project	21
6.	Related Works	22
7.	Conclusion	23
	Bibliography	24
A.	Appendix	25
A.1.	First Appendix Section	25

List of Figures

1.1. SDQ logo	3
A.1. A figure	25

List of Tables

1.1. A table 4

file : ///C : /Users/Burmi/AppData/Local/Temp/ForgetGate.pdf
https : //alr.anthropomatik.kit.edu/karrierebibel.de/curriculum-vitae/https : //www.indeed.com/career-advice/resumes-cover-letters/cv-format-guide

1. Motivation

1.1. Videogames and AI Research

1.1.1. Real-Time Strategy Games for AI

Partially Observable Environment (Fog of War), Complex interactions between Units (Micro Dynamic), Concurrent Actions, with multiple Agents on multiple scales (scale of one unit and one game participant).

Torchcraft: (insert bib link here)

- Extremely complex number of combinations (unit states, uncertainty (scouting) and volatile environment states) means that classic planning and search are not practical.
- partial observable environment, hard to quantify performance standard.
- makes a good benchmark for AI.

A Review of Real-Time Strategy Game AI:

- Video Games are an attractive alternative to robotics, as they provide a complex and realistic environment for simulations without the cost of real-world equipment and problems of practicality.

1.2. Multiagent Systems

Use Cases:

- Multiagent systems can be used in game theory and financing
- Reconnaissance robots covering a wide area. Communication not always possible.
- Smart Grid for Electricity, Power allocation
- Flow Line Systems
- Stock markets
- Load Balancing.
- Network Systems (IoT).
- Traffic Light Control
- Autonomous Driving
- Automating turbulence modelling (aircraft design, weather forecasting, climate prediction).

Aspects:

- Ant-Colony-Optimization, which can be used for learning.
- Emergent Behavior.
- Swarm Intelligence.
- multi-agent reinforcement learning.
- multi-agent learning.



Figure 1.1.: SDQ logo

- game theory
- compare these approaches to classical AI approaches in Videogames and RTS's (FSM, B-Trees, GOAP)

This is the SDQ thesis template. For more information on the formatting of theses at SDQ, please refer to <https://sdqweb.ipd.kit.edu/wiki/Ausarbeitungshinweise> or to your advisor.

1.3. Spacing and indentation

To separate parts of text in \LaTeX , please use two line breaks. They will then be set with correct indentation. Do *not* use:

- `\\`
- `\parskip`
- `\vskip`

or other commands to manually insert spaces, since they break the layout of this template.

1.4. Example: Citation

A citation: [1]

1.5. Example: Figures

A reference: The SDQ logo is displayed in Figure 1.1. (Use `\autoref{}` for easy referencing.)

1.6. Example: Tables

The `booktabs` package offers nicely typeset tables, as in Table 1.1.

1.7. Example: Formula

One of the nice things about the Linux Libertine font is that it comes with a math mode package.

$$f(x) = \Omega(g(x)) \ (x \rightarrow \infty) \Leftrightarrow \limsup_{x \rightarrow \infty} \left| \frac{f(x)}{g(x)} \right| > 0$$

abc	def
ghi	jkl
123	456
789	0AB

Table 1.1.: A table

2. Information to sort

2.1. Books

AI Game Programming Wisdom (1-4)

2.2. MARL - A Comprehensive Survey of Multiagent Reinforcement Learning - 2008

MARL - A Comprehensive Survey of Multiagent Reinforcement Learning - 2008 Benefits

- can be parallelized.
- can use experience sharing via communication, or with a teacher-learner relationship.
- Failure of one agent can be covered by other agents.
- insertion of new agents => scaleable.
- MARL Complexity: Exponential in number of agents.
- exploration (new knowledge) - exploitation (current knowledge) - Tradeoff.
- They explore about the environment and other agents.
- need for coordination.

Application Domains

- simulation better than real-life (better scalability and robustness).
- Distributed Control: for controlling processes (for larger industry plants).
 - avenue for future work.
 - used for traffic, power or sensory networks.
 - could also be used for pendulum systems.
- Robotic Teams (Multirobot):
 - simulated 2D space.
 - navigation: Reach a goal with obstacles. Area sweeping (retrieval of objects (also cooperative)).
 - pursuit: Capture a prey robot.
- Automated Trading: Exchange goods on electronic markets with negotiation and auctions.
- Resource Management: Cooperative team to manage resources or as clients. (routing, load balancing).

Practicality and Future works

- Scalability Problem: Q-functions do not scale well with the size of the state-action space.
 - Approximation needed: for discrete large state-action spaces, for continuous states and discrete actions or continuous state and action.
 - Heuristic in nature and only work in a narrow set of problems.
 - Could use theoretical results on single-agent approximate RL.

- also could use discovery and exploitation of the decentralized, modular structure of the multiagent task.
- MARL without prior knowledge is very slow.
 - Need humans to teach the agent.
 - shaping: first simple task then scale them.
 - could use reflex behavior.
 - Knowledge about the task structure.
- Incomplete, uncertain state measurements could be handled with partial observability techniques (Markov decision process).
- Multiagent Goals needs a stable learning process for the environment and an adaption for the dynamics of other agents.
- using game-theory-based analysis to apply to the dynamics of the environment.

2.3. Artificial Intelligence - A modern Approach

2.3.1. Agents and Environments

p.34

- **agent**: anything that perceives its **environment** through **sensors** and acting upon that environment using **actuators**.
- **percept**: agent's perceptual inputs at any given instance. Percept sequence is a complete history of perception.
- agents choice of action decided upon the history of perception, but not anything it has not perceived.
- its behavior is described by the **agent function**, which is internally implemented by the **agent program**.

2.3.2. Rational Agent

p.36

- **rational agent**: it does the correct thing. Correctness is determined by a performance measure, which is determined by the changed environment states.
- design **performance measures** according to what one actually wants in the environment, rather than according to how one thinks the agent should behave.
- rational depends on:
 - the performance measure that defines the criterion of success
 - the agent's prior knowledge of the environment.
 - The actions that the agent can perform.
 - The agent's percept sequence of data.
- depending on the measures the agent might be rational or not.
- an **omniscient agent** knows the actual outcome of its actions and can act accordingly, but this is impossible in reality.

- rationality maximizes expected performance, while perfection (omniscient) maximizes actual performance.
- agents can do actions in order to modify future percepts, called **information gathering, or exploration**.
- rational agents learn as much as possible from what it perceives.
- his knowledge can be augmented and modified as it gains experience.
- if the agent relies on the prior knowledge of its designer rather than on its own percepts, we say that the agent lacks **autonomy**.
- it should learn what it can to compensate for partial or incorrect prior knowledge.
- give it some initial knowledge and the ability to learn, so it will become independent of its prior knowledge.

2.3.3. Nature of Environments

p.40

- **task environments**: the “problems” to which rational agents are the “solutions”.
- Describe the task environment in the following aspects P(Performance measure), E(Environment), A(Actuators), S(Sensors).
- **fully observable**: the agent’s sensors give it access to the complete state of the environment. All aspects that are relevant to the choice of actions
- **partially observable**: otherwise. Because of missing sensors or noise.
- no sensors: unobservable
- single-agent environments and multi-agent environments.
- multi-agent can be either competitive (chess) or cooperative (avoiding collisions maximizes performance).
- **communication** emerges as a rational behavior in multiagent environments.
- randomized behavior is rational because it avoids the pitfalls of predictability.
- **Deterministic**: next state of environment is completely determined by the current state and the action executed by the agent, otherwise it is **stochastic**.
- you can ignore uncertainty that arises purely from the actions of other agents in a multiagent environment.
- If the environment is partial observable, it could appear to be stochastic, which implies quantifiable outcomes in terms of probabilities.
- an environment is **uncertain** if it is not fully observable or not deterministic.
- **episodic**: the agent’s experience is divided into atomic episodes. In each the agent receives a percept and performs a single episode. The next episode does not depend on the actions taken in previous episodes, otherwise it is **sequential**.
- When the environment can change while the agent is deliberating, then the environment is **dynamic** for that agent otherwise it is **static**.
- if the environment itself does not change with the passage of time but the agent’s performance score does, then we say the environment is **semi dynamic**.
- **discrete/continuous** applies to the state of the environment, to the way time is handled, and to the percepts and actions of the agents.
- **known vs. unknown**: refers to the agent’s state of knowledge about the “laws of physics” of the environment. Known environment, the outcomes for all actions are given,

otherwise the agent needs to learn how it works. An environment can be known, but partially observable (solitaire: I know the rules but still unable to see the cards that have not yet been turned over)

- hardest case: partially observable, multiagent, stochastic, sequential, dynamic, continuous, and unknown
- **environment class**: multiple environment scenarios to train it for multiple situations.
- you can create an **environment generator**, that selects environments in which to run the agent.

2.3.4. Structure of Agents

p.46

- agent = architecture (computing device) + program (agent program).
- agent programs take the current percept as input and return an action to the actuators.
- agent program takes the current percept, agent function which takes the entire percept history.
- **table driven agent**: Uses a table of actions indexed by percept sequences. This table grows way to fast and is therefore not practical.

Simple reflex agents:

- **simple reflex agents**: Select the actions on the basis of the current percept, ignoring the rest of the history.
- **condition-action-rule**: these agents create actions in a specific condition (if-then). These connections can be seen as reflexes.
- uses an **interpret-input** function as well as a **rule-match** function.
- they need the environment to be fully observable. They could run into infinite loops.
- you can mitigate this by using randomization for the actions. Which is non-rational for single agent environments.

Model-based reflex agents:

- keep track of the part of the world an agent cannot see now. It maintains some sort of **internal state** that depends on the percept history.
- agents need to know how the world evolves independently of the agent and how the agent's own actions affect the world.
- with this it creates a **model** of the world hence it is called model-based agent.
- it needs to update this state given sensor data.
- this model is a **best guess** and does not determine the entire current state of the environment exactly.

Goal-based agents:

- an agent needs some sort of **goal information** that describes situations that are desirable. This can also be combined with the model.
- Usually agents need to do multiple actions to fulfill a goal which requires **search** and **planning**.
- this also involves consideration of the future.
- the goal-based agent's behavior can be easily changed to go to a different destination by using a goal where a reflex agent needs completely new rules.

Utility-based agents:

- goals provide a crude binary distinction between good and bad states.
- use an internal **utility function** to create a performance measure.
- if the external performance measure and the internal utility function agree, the agent will act rationally.
- if you have conflicting goals the utility function can specify the appropriate **tradeoff**.
- if multiple goals cannot be achieved with certainty, utility provides a way to determine the **likelihood** of success.
- a rational utility-based agent chooses the action that **maximizes the expected utility**.
- any rational agent must behave as if it possesses a utility function whose expected value it tries to maximize.
- a utility-based agent must model and keep track of its environment.

Learning Agents:

- it allows the agent to operate in initially unknown environments and to become more competent than its initial knowledge alone might allow.
- 4 conceptual components: **learning element** (responsible for improvements), **performance element** (select external action), **critic** (gives feedback to change the learning element), **problem generator** (suggesting actions that lead to new and informative experiences).
- critic tells the learning element how well the agent is doing given a performance standard. It tells the agent which percepts are good and which are bad.
- problem generator allows for exploration and suboptimal actions to discover better actions in the long run.
- learning element: simplest case: learning directly from the percept sequence.
- the **performance standard** distinguishes part of the incoming percept as a reward or penalty that provides direct feedback on the quality of the agent's behavior.

How the components of agent programs work:

- **atomic representation**: Each state of the world is indivisible. Algorithms like search and game-playing, Hidden Markov models and Markov decision models work like this.
- **factored representation**: splits up each state of a fixed set of variables or attributes which each can have a value. Used in constraint satisfaction algorithms, propositional logic, planning, Bayesian networks.
- **structured representation**: here the different states have connections to each other. Used in relational databases, first-order logic, first-order probability models, knowledge-based learning and natural language understanding.
- more complex representations are more **expressive** and can capture everything more concise.

2.3.5. Multiagent Planning

p.425

- each agent tries to achieve its own goals with the help or hindrance of others
- wide degree of problems with various degrees of **decomposition of the monolithic agent**.

- multiple concurrent effectors => **multieffector planning** (like type and speaking at the same time).
- effectors are physically decoupled => **multibody planning**.
- if relevant sensor information for each body can be pooled centrally or in each body like single-agent problem.
- When communication constraint does not allow that: **decentralized planning problem**. planning phase is centralized, but execution phase is at least partially decoupled.
- single entity is doing the planning: one goal, that every body shares.
- When bodies do their own planning, they may share identical goals.
- **multibody**: centralized planning and execution send to each.
- **multiagent**: decentralized local planning, with coordination needed so they do not do the same thing.
- Usage of **incentives** (like salaries) so that goals of the central-planner and the individual align.

Multiple simultaneous actions:

- **correct plan**: if executed by the actors, achieves the goal. Though multiagent might not agree to execute any particular plan.
- **joint action**: An Action for each actor defined => joint planning problem with branching factor b^n (b = number of choices).
- if the actors are **loosely coupled** you can describe the system so that the problem complexity only scales linearly.
- standard approach: pretend the problems are completely decoupled and then fix up the interactions.
- **concurrent action list**: which actions must or most not be executed concurrently. (only one at a time)

Multiple agents: cooperation and coordination

- each agent makes its own plan. Assume goals and knowledge base are shared.
- They **might choose different plans** and therefore collectively not achieve the common goal.
- **convention**: A constraint on the selection of joint plans. (cars: do not collide is achieved by “stay on the right side of the road”).
- widespread conventions: social laws.
- absence of convention: use communication to achieve common knowledge of a feasible joint plan.
- The agents can try to **recognize the plan other agents want to execute** and therefore use plan recognition to find the correct plan. This only works if it is unambiguously.
- an **ant** chooses its role according to the local conditions it observes.
- ants have a convention on the importance of roles.
- ants have some learning mechanism: a colony learns to make more successful and prudent actions over the course of its decades-long life, even though individual ants live only about a year.
- Another Example: **Boid**
- If all the boids execute their policies, the flock inhibits the emergent behavior of flying as a pseudorigid body with roughly constant density that does not disperse over time.

- **most difficult multiagent** problems involve both cooperation with members of one's own team and competition against members of opposing teams, all without centralized control.

2.3.6. Game Theory

p.666

2.3.7. Mechanism Design for Multiple Agents

p.679

2.3.8. Adversarial Search

p.182

2.3.9. Probabilistic Reasoning over Time

p.587

2.3.10. Reinforcement Learning

p.830

2.3.11. Planning Uncertain Movements (Potential Fields)

p.993

2.4. RTSAI - A Review of Real-Time Strategy Game AI - 2014

2.4.1. Tactical Decision making

- In Practice: FSM mostly used.
- Tactical and micromanagement decisions can be separated from other aspects of the game.
- learning domain knowledge is feasible
- Less focus in research.

Reinforcement Learning

- Reinforcement Learning can be used.
- requires clever state abstraction to learn effectively.
- RL not used for Strategic aspects as the problem space is larger and the reward function has delays.
- Used to learn an expected reward value for a particular actions.

- Using a hierarchical approach, where a unit is controlled individually with a higher-level group control affecting each individual's decision would decrease the complexity and make approaches more viable.

GameTree Search

- Much Simplified scenario with a known deck of strategies and counters.
- Only then can a search algorithm work with skipping uninteresting states and the moves only be described with pairs and counters. (Nash).
- alpha-beta search / pruning.
- These simplification do not work well into a real game because of the lack of unit collisions and accelerations. Really bad at hit-and-run maneuver.
- For StarCraft: The Search will only allow up to eight units per side in a 2player battle before it is too slow.

Monte Carlo Planning

- sample decision space using randomly generated plans to find successful ones.
- Monte Carlo tree search may be effective at strategic level decision making for RTS Games.
- Using upper confidence bounds applied to trees (UCT). Plays game randomly out to a terminal state.

Other Models

- using objectives, opportunities and threats in a bayesian model learned with RL. Paired with hierarchical FSM for different behavior sets for different scenarios.
- Case-Based Reasoning: selects most similar case for reuse, that only is executed when it is able to be executed. Also used reactionary cases (reflexes)
- Neuroevolution: Evolutionary Algorithms for Neural Networks: Each Unit has its own neural network. (rtNEAT). Environmental sources and actions as outputs. How well does it scale for more units?

2.4.2. Strategic Decision making

- In Research: Planning Systems. They determine a sequence of actions to be taken.
- Challenging, because of incomplete information.
- Goal: Humanlike level of competence and reduce development effort for the industry for scripting methods.
- Main Techniques used: case-based planning (CBP), goal-drive autonomy (GDA) and hierarchical planning.
- Kohan 2: Kings of War: Fun AI with Goals and priorities that resulted in a fun opponent. Also easy to update.

Case-Based Planning

- finds similar past situations from which to draw potential solutions to the current situation.
- for CBP: Solutions are a set of potential plans or subplans that are likely to be effective.
- They exhibit poor reactivity at the strategic level and excessive reactivity at the action level.
- Can abstract the states into a state lattice of possible orders (build orders) combined with small set of features and abstracted actions.
- Approaches with using game logs and behaviors, goals and alive conditions.

- During plan execution it may be modified in order to adapt to unforeseen events. Examples use decision tree models. Decision tree predicts the high-level situation which determines the attributes and their weights for the case selection. The Tree is learnt.
- can be combined with fuzzy sets to abstract state information and abstracting the regions of the map and their state with choke points, military presence, combat intensity, lost units and amounts of each friendly and enemy unit type.
- needs replays (some examples could achieve a 60% winrate after one replay)

hierarchical Planning

- breaking up the problem hierarchically reduces the complexity of the problem, but create new issues in coordination between the different levels
- hierarchical plans maps well to the hierarchy of goals and subgoals in typical RTS games.
- Hierarchical Task Networks (HTN): contains tasks, their ordering and methods for achieving them. high level tasks can be decomposed up until concrete actions.
- They allow the agent to react dynamically to problems.
- There has not been much work in learning HTNs for RTS only for simpler domains.
- can also use an active behavior tree which has parallel, sequential and conditional behaviors and goals in a tree structure. Here the tree is expanded during the execution by selecting behaviors.
- Can combine this with other methods like hierarchical CBP.

Behavior Trees

- a hierarchy of decision and action nodes, used a lot by game designers.
- they can be edited using visual tools, which make them more accessible to non-programmers.
- as they are hierarchical they can cover a wide range of behavior.
- They have been used to enable direct control of a case-based-planner's behavior: Combining Expert Knowledge and Learning from Demonstration in Real-Time Strategy Games. (seems nice)

Goal-Driven Autonomy

- An agent reasons about its goals, identifies when they need to be updated and changes or adds to them as needed for subsequent planning and execution.
- Addresses the high- and low-level reactivity problem experienced by CBP by reasoning about and reacting to why a goal is succeeding or failing.
- unexpected situation: record it as a discrepancy, generate an explanation and form new goals to revise the plan.
- There is a CBR and RL Version of GDA that allows it to learn new domain knowledge which is not needed to be fed to it.

2.4.3. Open Research Areas

Game AI in Industry

- In Practice: Industry uses older and simpler approaches.
- papers reason that AI Research will be useful for development by reducing the work involved in creating AI opponents, or by allowing game developers to create better AI opponents.
- Planetary Annihilations uses flow fields for path-finding and neural networks for squad unit control.

- industry tries to create challenging but defeatable opponents that are fun to play against.
- the game is more fun when reasonably challenging.
- the agent tries to win at all cost, but the game industry wants an agent that loses in a more human-like way. it is fun to finding and taking advantage of opportunities and opponent's mistakes.
- usually classic AI techniques are enough to make a fun AI. It can cheat.
- the old techniques are predictable, reliable and easy to test and debug. Academic AI are seen as difficult to customize, tune, or tweak.
- could gain a lot by making the AI more Humanlike: capable of cooperation, using surprises, deception, distractions and coordinated attacks, planning effective strategies and changing strategies to become less predictable. (Scott 2002)
- Being predictable and exploitable in the same fashion over multiple games: not fun.
- AI can make mistakes: They need to seem plausible in the context of the game. Something a human would do.
- Replicate Human Playstyle (Turtle-ing), which gives the AI more personality. This could be trained and learned from a human player (from demonstration and offline). Copy how player play the game.
- More work to be done how to make complex systems easier to tweak and customize. Produce specific behavior while retaining learning or reasoning capabilities.

Multiscale AI

- currently bots require multiple abstractions and reasoning mechanics to work in concert. They usually separate ways of handling tactical and strategic level decision making. Separately managing resources, construction and reconnaissance.
- each module interacts with each other, so straightforward hierarchy of command is difficult, as high level systems require direct knowledge of other systems.
- This is called multiscale AI problems.
- Some approaches use high-level commanders and sub-commanders. Commanders further down the hierarchy are increasingly focused on a particular task, but have less information about the overall game state, so therefore must rely on their parents to make them act appropriately in the bigger picture. They can use information from their parent in the hierarchy, but they are unable to react and coordinate with other low-level controllers directly for cooperative actions.

2.4.4. AI Techniques that can be used

2.5. Ant Colony Optimization

2.5.1. Wikipedia Article

Ant Colony Optimization Algorithm, Wikipedia

- is used for solving computational problems which can be reduced to finding good paths through graphs.
- artificial ants locate optimal solutions by moving through a parameter space representing all possible solutions.

Tactical Decision Making	Strategic Decision Making and Plan Recognition
Reinforcement Learning	Case-Based Planning
Game-Tree Search	Hierarchical Planning
Bayesian Models	Behavior Trees
Case-Based Reasoning	Goal-Driven Autonomy
Neural Networks	State Space Planning
	Evolutionary Algorithms
	Cognitive Architectures
	Deductive Reasoning
	Probabilistic Reasoning
	Case-Based Reasoning

- they record their positions and the quality of their solutions for later iterations to find better solutions (pheromones).

2.6. UNSORTED

Gordon 2000: Ants at Work.

Gordon 2007: Control without hierarchy. Nature.

Links:

Ant Simulation Video 1

Ant Simulation Video 2

Boids Video

Distributed Artificial Intelligence, Wikipedia

Multi-agent learning, Wikipedia

Bees algorithm, Wikipedia

Swarm Intelligence, Wikipedia

2.7. References & Papers

2.7.1. Ant Colony Optimization (ACO)

ACO - Ant Colony Optimization for learning Bayesian network - 2002

2.7.2. Multi Agent Reinforcement Learning (MARL)

MARL - Multiagent Reinforcement Learning - Theoretical Framework and an Algorithm - 1998

MARL - Deep Reinforcement Learning for Robot Swarms - 2019 - KIT

MARL - Hierarchical multi-agent reinforcement learning - 2006

MARL - Learning to Communicate with Deep Multi-Agent Reinforcement Learning - 2016

MARL - Reward shaping for knowledge-based multi-objective multi-agent reinforcement learning - 2017

MARL - GAMA - Graph Attention Multi-agent reinforcement learning algorithm for cooperation - 2020
MARL - Plan-based reward shaping for multi-agent reinforcement learning - 2016
MARL - Multi-Agent Reinforcement Learning Using Linear Fuzzy Model Applied to Cooperative Mobile Robots - 2018
MARL - Multi-Agent Reinforcement Learning - a critical survey - 2003
MARL - Stabilising Experience Replay for Deep Multi-Agent Reinforcement Learning - 2017
MARL - Mean Field Multi-Agent Reinforcement Learning - 2018
MARL - Multi-Agent Reinforcement Learning A Selective Overview of Theories and Algorithms - 2021
MARL - Coordinating multi-agent reinforcement learning with limited communication - 2013
MARL - Multi-Agent Reinforcement Learning A Report on Challenges and Approaches - 2018

2.7.3. Multiagent Systems (MAS)

MAS - Multi-Agent Systems - A Survey - 2018
MAS - Distributed Cooperative Control and Communication for Multi-agent Systems - 2021
MAS - PRIMA 2020 Principles and Practice of Multi-Agent Systems - 2021
MAS - The Multiagent Planning Problem - 2016
MAS - Swarm Intelligence - 2010
MAS - Swarm Intelligence - 2012
MAS - Swarm Intelligence - 2014
MAS - Swarm Intelligence - 2016
MAS - Swarm Intelligence - 2018
MAS - Swarm Intelligence - 2020
MAS - Using Multi-Agent Potential Fields in Real-Time Strategy Games - 2008
MAS - Evaluating the Effectiveness of Multi-Agent Organisational Paradigms in a Real-Time Strategy Environment, Engineering Multiagent Systems Track - 2019
MAS - The StarCraft Multi-Agent Challenge - 2019
MAS - Neuroevolution based multi-agent system for micromanagement in real-time strategy games - 2012
MAS - Dealing with fog of war in a Real Time Strategy game environment - 2008

- p.3: The Initial set of charges were found using trial and error => can use learning for that.
- using multi-agent potential fields for unit movements.

2.7.4. RTS AI

RTSAI - A Review of Real-Time Strategy Game AI - 2014
RTSAI - Artificial Intelligence Techniques on Real-time Strategy Games - 2018
RTSAI - Informed Monte Carlo Tree Search for Real-Time Strategy games - 2016
RTSAI - Case based Reasoning Research (Palmer) - 2011 - CBP - BT

2.7.5. RTS Learning

RTSLearn - Combining AI Methods for Learning Bots in a Real-Time Strategy Game - 2008

RTSLearn - Combining Strategic Learning and Tactical Search in RTS Games - 2017

RTSLearn - Imitative Learning for RTS - 2012

RTSLearn - Deep RTS A Game Environment for Deep Reinforcement Learning in Real-Time Strategy Games - 2018

RTSLearn - Learning Map-Independent Evaluation Functions for Real-Time Strategy Games - 2018

RTSLearn - Learning Probabilistic Behavior Models in Real-Time Strategy Games - 2011

RTSLearn - Machine Learning and Games - 2006

RTSLearn - Real Time Strategy Games: A Reinforcement Learning Approach - 2015

RTSLearn - Torchcraft - A Library for Machine Learning Research on Real-Time Strategy Games - 2016

3. Fundamentals

Topics:

- multiagent/multibody Systems (MAS).
 - MAS Reinforcement Learning
 - * They use stochastic games (Markov Games) as generalization of Markov Decision Processes.
 - Hierarchical MAS, Hierarchical Reinforcement Learning for MAS.
 - MAS with Cooperation and Competition.
 - Particle Swarms (nicht so meins).
 - Problems:
 - * MAS Movement Problems (Potential Fields).
 - * ? using MAS for Moving a Multi-Legged Robot (Spider-like) with a navigation problem design as a hierarchical MAS?
 - * Path Planning Navigation with Heterogeneous Agents?
 - * MAS Task Problems: Rendezvous, Pursuit Evasion (Single and one Evader) (Boid?), (MAS - Deep Reinforcement Learning for Robot Swarms - 2019 - KIT)
 - * Multi-Agent Path Finding (MAPF). Scalability for this: For fixed space they get into each others way.
 - * Collective Foraging: (Ants-kind). Problem when communication only happens in an area, use local information exchange groups. Information Transfer. (MAS - Swarm Intelligence - 2020), Preferential Foraging (MAS - Swarm Intelligence - 2018 - p.289)
 - * Coverage: Multi-robot Information Gathering / Scouting. (MAS - Swarm Intelligence - 2020), Pattern Formation (MAS - Swarm Intelligence - 2016 - p.14)
 - * Coalition: Heterogeneous Group of Agents. They have different skills / attributes that affect the environment. Like an Ant Caste System? Some Agents have better sensors? Only some agents have some sensors? What if the specialization is taken to the extreme? (MAS - Swarm Intelligence - 2020). Limited Visibility Sensors (MAS - Swarm Intelligence - 2018 - p.56). Going from a homogeneous group of agents, to randomized specialities, to extrem specializations. How does it change? Mixed with an Hierarchical Approach? They need to find groups to work together? Some are fast (but cannot see much, there is an insect that cannot see while running), Some have good sensors. Communication range? Genetic Diversity, Task-Allocation and Task-Switching (MAS - Swarm Intelligence - 2016 - p.109)
 - * Collective Gradient Perception: Using Abilities of other Agents to take advantage of the whole group. (Flocking) (MAS - Swarm Intelligence - 2020)
 - * Indirect Communication through changing states in the environment (birds transport something via cable). Also like using Pheromone Trails (Quality-

- Sensitive Foraging through virtual pheromone trails). (MAS - Swarm Intelligence - 2018 - p.15 - p.147)
- * Control Architecture: Behavior Trees, FSM. (MAS - Swarm Intelligence - 2018 - p.42)
- * Maze-Like Environment with Ant Algorithms (MAS - Swarm Intelligence - 2018 - p.162)
- * Search and Rescue? (Kinda like Foraging?)
- * Disruption: Disrupting Aspects of the Swarm and how they react to it, Swarm Attack: (MAS - Swarm Intelligence - 2018 - p.225), Coherence of Collective Decision Making (MAS - Swarm Intelligence - 2018 - p.264)
- * Evolutionary Systems: NEAT (MAS - Swarm Intelligence - 2012 - p.98)
- Graph-Based Visualisation for MAS. (MAS - Swarm Intelligence - 2010). How do you visualize them?
- Adding new Agents to the mix, can you use transfer learning?
- Adaptive Learning for MAS?
- Control System: Either fully self-organizing or completely centralized. Hybrid Control of Swarms (MAS - Swarm Intelligence - 2018 - p.69)
- Simulation of MAS: ARGoS
- Best-of-n Problem: Swarm selects best option out of n alternatives. (MAS - Swarm Intelligence - 2018 - p.251)
- Sensory Errors for Foraging, Dynamic Task Partitioning (MAS - Swarm Intelligence - 2016 - p.124), Task Partitioning Problem (MAS - Swarm Intelligence - 2012 - p.122)
- Task Hierarchy, Multi-Objective
- Random Walks as a search strategy (MAS - Swarm Intelligence - 2016 - p.196)
- Ant-Colony-Optimization (ACO)
 - ACO for learning.
 - ACO for classification: (MAS - Swarm Intelligence - 2016 - p.64)
- General RTS AI Approaches.
 - FSM
 - Behavior Trees
 - GOAP
- RTS-Learning
 - Reinforcement Learning
 - Learning Probabilistic Models
 - CNN

4. Problem and Approaches

4.1. Definition of the Problem domain

Agent:

- performance measure: which one?
- needs information gathering
- uses percepts to find correct action from prior knowledge
- communication is important
- learning while a game is running?

Multiagent:

- multiagent (decentralized planning with coordination) or multibody (centralized planning)?
- multi-effector (units can walk and attack) (can be made easier).

Environment:

- partially observable (fog of war).
- cooperative multiagent environment for a given player.
- competitive multiagent environment between players.
- usually deterministic, games may use RNG than it would be stochastic.
- stochastic because of partial observability.
- sequential (current actions depend on previous actions).
- static environment (when AI is frame-locked), otherwise dynamic
- pseudo-continuous environment
- known, but partially observable.
- basically hardest case: partially observable, multiagent, stochastic, sequential, dynamic, continuous and unknown.

4.2. Problems and Aspects for an RTS AI

- balancing changes.

4.2.1. Tactical Decision Making

Tactics. can use some metric (max enemy firepower removed in shortest time). Video game industry uses FSMs. Exploration approaches here possible.

4.2.2. Strategic Decision Making

5. Project

- use a video game environment to simulate ants and learning with multi-agent systems.
- could create a survival scenario (vs. nature) and an adversarial scenario (vs. another AI or player).
- create abstracted mechanics that create a learning environment in the context of a video game.
- Aspects of the AI: tactical (micro) and strategic (macro) decision making, plan recognition and learning.

6. Related Works

Is this part of the motivation???

7. Conclusion

Bibliography

- [1] Steffen Becker, Heiko Koziolk, and Ralf Reussner. “The Palladio Component Model for Model-driven Performance Prediction”. In: *Journal of Systems and Software* 82 (2009), pp. 3–22. DOI: 10.1016/j.jss.2008.03.066. URL: <http://dx.doi.org/10.1016/j.jss.2008.03.066>.

A. Appendix

A.1. First Appendix Section

Figure A.1.: A figure

...